

Introduction to R

Dr. Arabin Kumar Dey

Assistant Professor
Department of Mathematics
Indian Institute of Technology Guwahati

December 28, 2011

Outline

- 1 Installing R
- 2 Why using R ?
- 3 Working with Data Set
- 4 Using Data Available in R

Outline

- 1 Installing R
- 2 Why using R ?
- 3 Working with Data Set
- 4 Using Data Available in R

Installing R on Linux/Windows

- Go to <http://cran.r-project.org>
- Select to download the latest version:
- Install and Open. The R window should look like :

Installing R

Why using R ?

Working with Data Set

Using Data Available in R



R Console

R : Copyright 2006, The R Foundation
Version 2.3.1 (2006-06-01)
ISBN 3-900051-07-0

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> |

R GUI

File Edit Misc Packages Windows Help

R Console

R : Copyright 2004, The R Foundation
Version 2.0.0 (2004-10-04)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

tmux.1.bash

5

Group1 Group2

R Data Editor

DV	Treatment
1	Group1
2	Group1
3	Group1
4	Group2
5	Group2
6	Group2

Outline

- 1 Installing R
- 2 Why using R ?
- 3 Working with Data Set
- 4 Using Data Available in R

- Complete statistical package and programming language
- Efficient functions and data structures for data analysis
- Powerful graphics
- Access to fast growing number of analysis packages
- Technical advantages: free, open-source, available for all OSs

- Complete statistical package and programming language
- Efficient functions and data structures for data analysis
- Powerful graphics
- Access to fast growing number of analysis packages
- Technical advantages: free, open-source, available for all OSs

- Complete statistical package and programming language
- Efficient functions and data structures for data analysis
- Powerful graphics
- Access to fast growing number of analysis packages
- Technical advantages: free, open-source, available for all OSs

- Complete statistical package and programming language
- Efficient functions and data structures for data analysis
- Powerful graphics
- Access to fast growing number of analysis packages
- Technical advantages: free, open-source, available for all OSs

- Complete statistical package and programming language
- Efficient functions and data structures for data analysis
- Powerful graphics
- Access to fast growing number of analysis packages
- Technical advantages: free, open-source, available for all OSs

Startup/Closing Behavior

- Starting R The R GUI versions under Windows and Mac OS X can be opened by double-clicking their icons.
- Alternatively, one can start it by typing 'R' in a terminal (default under Linux).
- The R environment is controlled by hidden files in the startup directory: .RData, .Rhistory and .Rprofile (optional).
- ## Closing R
 > q()
 Save workspace image?
 [y/n/c] :
- Note When responding with 'y', then the entire R workspace will be written to the .RData file which can become very large. Often it is sufficient to just save an analysis protocol in an R source file. This way one can quickly regenerate all data sets and objects.

Startup/Closing Behavior

- Starting R The R GUI versions under Windows and Mac OS X can be opened by double-clicking their icons.
- Alternatively, one can start it by typing 'R' in a terminal (default under Linux).
- The R environment is controlled by hidden files in the startup directory: .RData, .Rhistory and .Rprofile (optional).
- ## Closing R
 > q()
 Save workspace image?
 [y/n/c] :
- Note When responding with 'y', then the entire R workspace will be written to the .RData file which can become very large. Often it is sufficient to just save an analysis protocol in an R source file. This way one can quickly regenerate all data sets and objects.

Startup/Closing Behavior

- Starting R The R GUI versions under Windows and Mac OS X can be opened by double-clicking their icons.
- Alternatively, one can start it by typing 'R' in a terminal (default under Linux).
- The R environment is controlled by hidden files in the startup directory: `.RData`, `.Rhistory` and `.Rprofile` (optional).
- `## Closing R`
`> q()`
Save workspace image?
`[y/n/c] :`
- Note When responding with 'y', then the entire R workspace will be written to the `.RData` file which can become very large. Often it is sufficient to just save an analysis protocol in an R source file. This way one can quickly regenerate all data sets and objects.

Startup/Closing Behavior

- Starting R The R GUI versions under Windows and Mac OS X can be opened by double-clicking their icons.
- Alternatively, one can start it by typing 'R' in a terminal (default under Linux).
- The R environment is controlled by hidden files in the startup directory: .RData, .Rhistory and .Rprofile (optional).
- ## Closing R
 > q()
 Save workspace image?
 [y/n/c] :
- Note When responding with 'y', then the entire R workspace will be written to the .RData file which can become very large. Often it is sufficient to just save an analysis protocol in an R source file. This way one can quickly regenerate all data sets and objects.

Startup/Closing Behavior

- Starting R The R GUI versions under Windows and Mac OS X can be opened by double-clicking their icons.
- Alternatively, one can start it by typing 'R' in a terminal (default under Linux).
- The R environment is controlled by hidden files in the startup directory: .RData, .Rhistory and .Rprofile (optional).
- ## Closing R
 > q()
 Save workspace image?
 [y/n/c] :
- Note When responding with 'y', then the entire R workspace will be written to the .RData file which can become very large. Often it is sufficient to just save an analysis protocol in an R source file. This way one can quickly regenerate all data sets and objects.

Installation of Cran package

- Linux - `install.packages("mypackage_1.0.tar.gz", repos=NULL)`
- Hands-on Experience ...

- ## Create an object with the assignment operator ' $<$ ' (or ' $=$ ')
> *object* < - ...
- ## List objects in current R session
> *ls()*
- ## Return content of current working directory
> *dir()*
- ## Return path of current working directory
> *getwd()*
- ## Change current working directory
> *setwd("/home/user")*

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my_script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my_script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my_script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Basic R syntax

- ## General R command syntax
 - > object < – function(arguments)
 - > object < – object[arguments]
- ## Execute an R script
 - > source("my script.R")
- ## Execute an R script from command-line
 - > R CMD BATCH my_script.R
 - > R –slave < my_script.R
- ## Finding help
 - > ?function
- ## Load a library
 - > library("my_library")
- ## Summary of all functions within a library
 - > library(help="my_library")

Data Type

- ## Numeric data: $1, 2, 3 > x < -c(1, 2, 3); x;$
`is.numeric(x); as.character(x)`
- ## Character data: `"a", "b", "c" > x < -c("1", "2", "3");`
`x; is.character(x); as.numeric(x)`
- ## Logical data: `TRUE, FALSE, TRUE`
`> x < -1 : 10 < 5; x`
`> !x`
- ## Return indices for the 'TRUEs' in logical vector
`> which(x)`

Data Type

- ## Numeric data: $1, 2, 3 > x < -c(1, 2, 3); x;$
`is.numeric(x); as.character(x)`
- ## Character data: `"a", "b", "c" > x < -c("1", "2", "3");`
`x; is.character(x); as.numeric(x)`
- ## Logical data: `TRUE, FALSE, TRUE`
`> x < -1 : 10 < 5; x`
`> !x`
- ## Return indices for the 'TRUEs' in logical vector
`> which(x)`

Data Type

- ## Numeric data: $1, 2, 3 > x < -c(1, 2, 3); x;$
`is.numeric(x); as.character(x)`
- ## Character data: $"a", "b", "c" > x < -c("1", "2", "3");$
`x; is.character(x); as.numeric(x)`
- ## Logical data: `TRUE, FALSE, TRUE`
 $> x < -1 : 10 < 5; x$
 $> !x$
- ## Return indices for the 'TRUEs' in logical vector
 $> which(x)$

Data Type

- ## Numeric data: $1, 2, 3 > x < -c(1, 2, 3); x;$
`is.numeric(x); as.character(x)`
- ## Character data: $"a", "b", "c" > x < -c("1", "2", "3");$
`x; is.character(x); as.numeric(x)`
- ## Logical data: `TRUE, FALSE, TRUE`
 $> x < -1 : 10 < 5; x$
 $> !x$
- ## Return indices for the 'TRUEs' in logical vector
 $> which(x)$

Caution !!

On the other hand, if you give a variable the same name as an existing function, R will treat the identifier as a variable if used as a variable, and will treat it as a function when it is used as a function:

```
c < -2 #typing c yields "2"
```

```
c(c, c) #yields a vector containing two 2s.
```

Better to avoid using 'c' as a variable

Data Objects

- ## Vectors (1D)
 `> myVec <- 1:10; names(myVec) <- letters[1:10]`
 `> myVec[1:5]; myVec[c(2,4,6,8)]; myVec[c("b", "d", "f")]`
- ## Matrices (2D), Data Frames (2D) and Arrays ($\geq 2D$)
 `> myMA <- matrix(1:30, 3, 10, byrow = T)`
 `> myDF <- data.frame(Col1=1:10, Col2=10:1)`
 `> myDF[1:4,]; myDF[,c("Col2", "Col1", "Col1")]`
- ## Lists: containers for any object type
 `> myL <- list(name="Fred", wife="Mary", no.children=3,
 child.ages=c(4,7,9))`
 `> myL[[4]][1:2]`
- ## Functions: piece of code
 `> myfct <- function(arg1,
 arg2, ...)`

General Subsetting Rules

- ## Subsetting by indices
 > myVec <- 1:26; names(myVec) <- LETTERS
 > myVec[1 : 4]
- ## Subsetting by same length logical vectors
 > myLog <- myVec > 10
 > myVec[myLog]
- ## Subsetting by field names
 > myVec[c(" B", " K", " M")]
- ## Special case
 > iris\$Species

Basic Operators and Calculations

- Comparison operators: $==$, $!=$, $<$, $>$, $<=$, $>=$
Example:
`> 1 == 1`
- Logical operators: *AND* : $\&$, *OR* : $|$, *NOT* : $!$
Example:
`> x < -1 : 10; y < -10 : 1`
`> x > y & x > 5`
- Calculations: ## Example:
`> x + y; sum(x); mean(x); sd(x); sqrt(x)`
`> apply(matrix(c(1,2,3,4,4,5,6,7,8,9,9,4),3,4), 1, mean)`

Outline

- 1 Installing R
- 2 Why using R ?
- 3 Working with Data Set**
- 4 Using Data Available in R

- ## Import Data into R
 > read.delim(" myData.csv") ??
 > a<-read.table(" myData.txt");
- Extract the portion you want to use, for example, $x < -a[, 2]$
- ## Export Data from R to File
 > write.table(myframe, file=" myfile.csv" , sep="\t" ,
 quote=F)

Some Great R Functions

- The **unique()** function to make vector entries unique
 `> unique(iris$Sepal.Length);`
 `length(unique(iris$Sepal.Length))`
- The **table()** function counts the occurrences of entries
 `> table(iris$Species)`
- The **aggregate()** function computes statistics of data aggregates
 `> aggregate(iris[,1:4], by=list(iris$Species), FUN=mean, na.rm=T)`
- The **%in%** function returns the intersect between two vectors
 `> month.name[month.name %in% c(" May", " July")]`
- The **merge()** function joins data frames based on a common key column
 `> merge(frame1, frame2, by.x=1, by.y=1, all = TRUE)`

Some Great R Functions

- The **unique()** function to make vector entries unique
 `> unique(iris$Sepal.Length);`
 `length(unique(iris$Sepal.Length))`
- The **table()** function counts the occurrences of entries
 `> table(iris$Species)`
- The **aggregate()** function computes statistics of data aggregates
 `> aggregate(iris[,1:4], by=list(iris$Species), FUN=mean, na.rm=T)`
- The **%in%** function returns the intersect between two vectors
 `> month.name[month.name %in% c(" May", " July")]`
- The **merge()** function joins data frames based on a common key column
 `> merge(frame1, frame2, by.x=1, by.y=1, all = TRUE)`

Some Great R Functions

- The **unique()** function to make vector entries unique
 `> unique(iris$Sepal.Length);`
 `length(unique(iris$Sepal.Length))`
- The **table()** function counts the occurrences of entries
 `> table(iris$Species)`
- The **aggregate()** function computes statistics of data aggregates
 `> aggregate(iris[,1:4], by=list(iris$Species), FUN=mean, na.rm=T)`
- The **%in%** function returns the intersect between two vectors
 `> month.name[month.name %in% c(" May", " July")]`
- The **merge()** function joins data frames based on a common key column
 `> merge(frame1, frame2, by.x=1, by.y=1, all = TRUE)`

Some Great R Functions

- The **unique()** function to make vector entries unique
 `> unique(iris$Sepal.Length);`
 `length(unique(iris$Sepal.Length))`
- The **table()** function counts the occurrences of entries
 `> table(iris$Species)`
- The **aggregate()** function computes statistics of data aggregates
 `> aggregate(iris[,1:4], by=list(iris$Species), FUN=mean, na.rm=T)`
- The **%in%** function returns the intersect between two vectors
 `> month.name[month.name %in% c(" May", " July")]`
- The **merge()** function joins data frames based on a common key column
 `> merge(frame1, frame2, by.x=1, by.y=1, all = TRUE)`

Some Great R Functions

- The **unique()** function to make vector entries unique
 `> unique(iris$Sepal.Length);`
 `length(unique(iris$Sepal.Length))`
- The **table()** function counts the occurrences of entries
 `> table(iris$Species)`
- The **aggregate()** function computes statistics of data aggregates
 `> aggregate(iris[,1:4], by=list(iris$Species), FUN=mean,`
 `na.rm=T)`
- The **%in%** function returns the intersect between two vectors
 `> month.name[month.name %in% c(" May", " July")]`
- The **merge()** function joins data frames based on a common key column
 `> merge(frame1, frame2, by.x=1, by.y=1, all = TRUE)`

Outline

- 1 Installing R
- 2 Why using R ?
- 3 Working with Data Set
- 4 Using Data Available in R

To use a data set available in one of the R packages, install that package (if needed). Load the package into R, using the `library()` function.

```
> library(MASS)
```

Extract the data set you want from that package, using the `data()` function. In our case, the data set is called `JohnsonJohnson`.

```
> data(JohnsonJohnson)
```

Working with Datasets in R

To use the variable names when working with data, use `attach()`:

```
> data(JohnsonJohnson)  
> attach(JohnsonJohnson)
```

After the variable names have been "attached", to see the variable names, use `names()`:

```
> names(JohnsonJohnson) To see the descriptions of the variables,  
use ?:  
> ?JohnsonJohnson
```

After modifying variables, use `detach()` and `attach()` to save the results:

```
# Make a copy of the data set  
johnson.copy <- JohnsonJohnson;  
detach(JohnsonJohnson)  
attach (johnson.copy)  
# Change the 10 th observation for JohnsonJohnson  
johnson.copy[10,1] <- 999
```

Caution!!

Avoid using `attach()` if possible. Many strange things can occur if you accidentally attach the same data frame multiple times, or forget to detach. Instead, you can refer to a variable using `$`.

Moral of the story

"attach at your own risk!"